

Introduction to time series and stationarity

FORECASTING USING ARIMA MODELS IN PYTHON



James Fulton

Climate informatics researcher

Motivation

Time series are everywhere

- Science
- Technology
- Business
- Finance
- Policy

Course content

You will learn

- Structure of ARIMA models
- How to fit ARIMA model
- How to optimize the model
- How to make forecasts
- How to calculate uncertainty in predictions

Loading and plotting

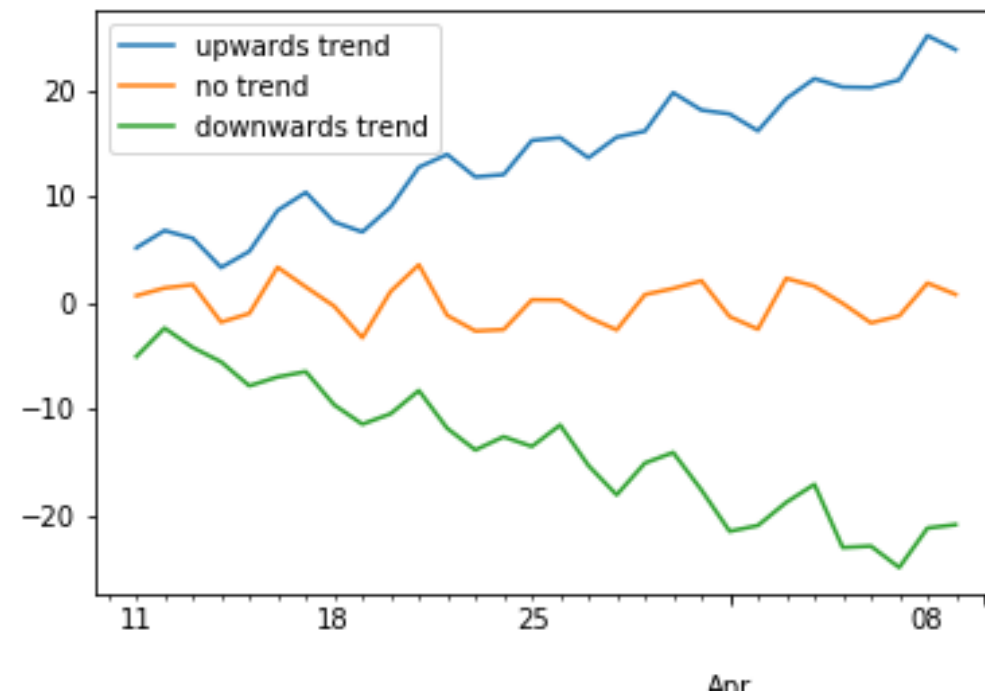
```
import pandas as pd
import matplotlib as plt

df = pd.read_csv('time_series.csv', index_col='date', parse_dates=True)
```

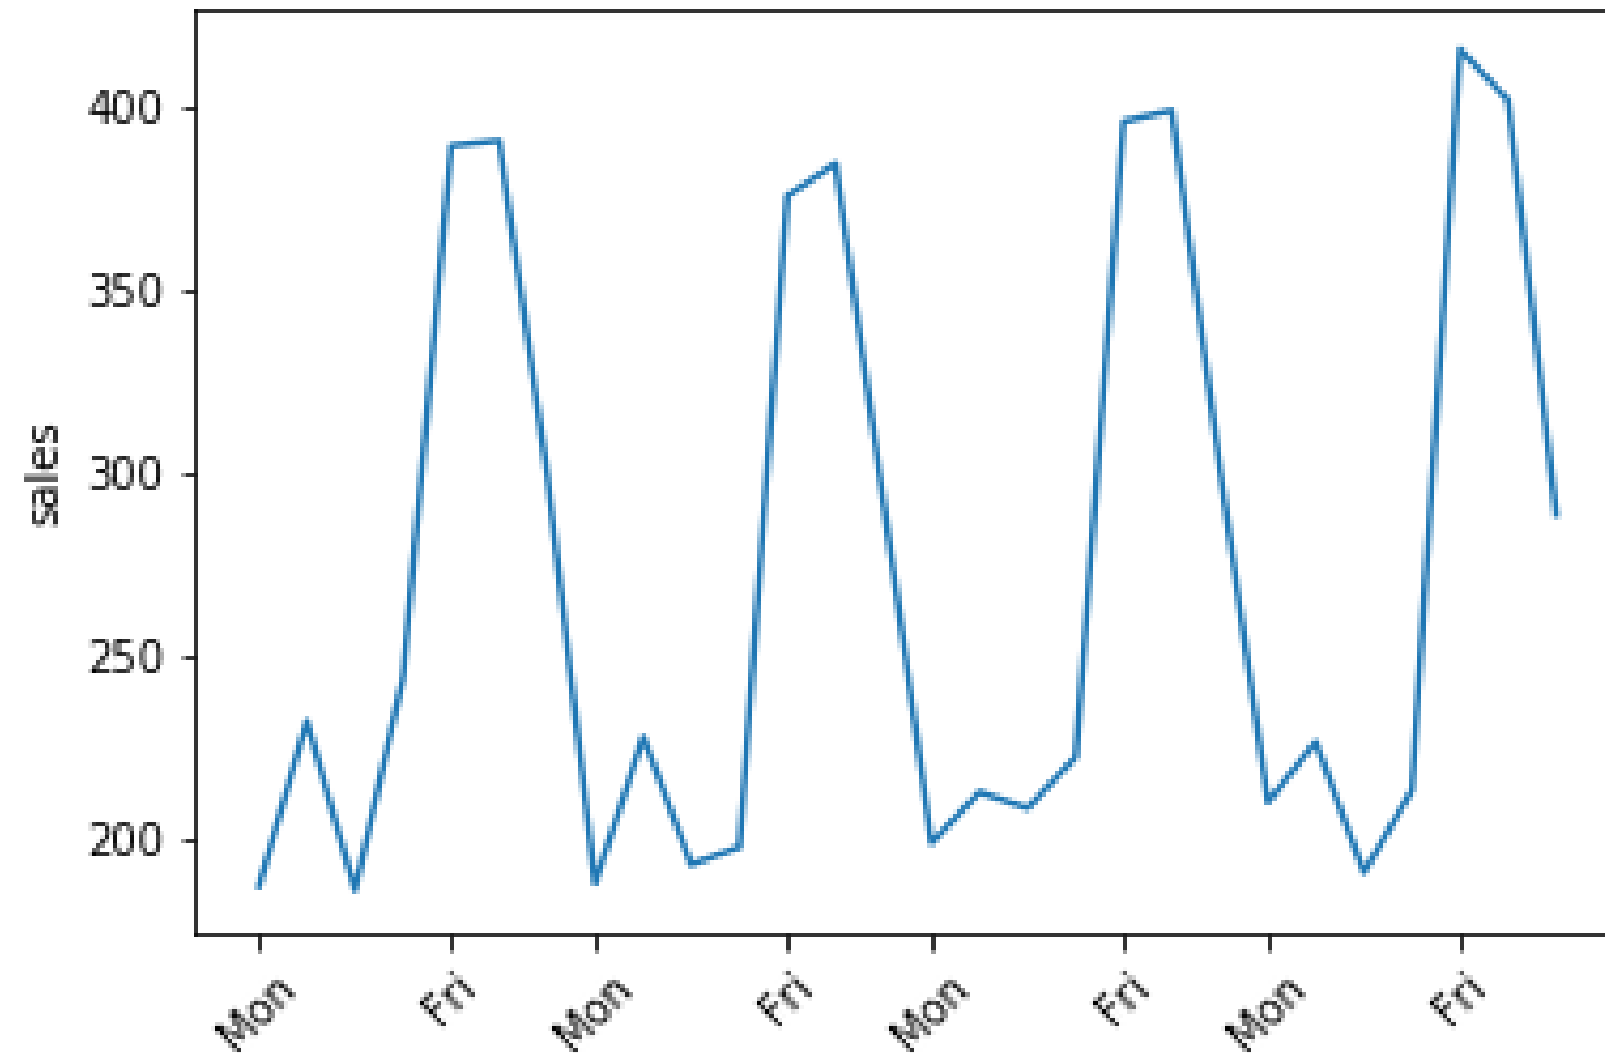
date	values
2019-03-11	5.734193
2019-03-12	6.288708
2019-03-13	5.205788
2019-03-14	3.176578

Trend

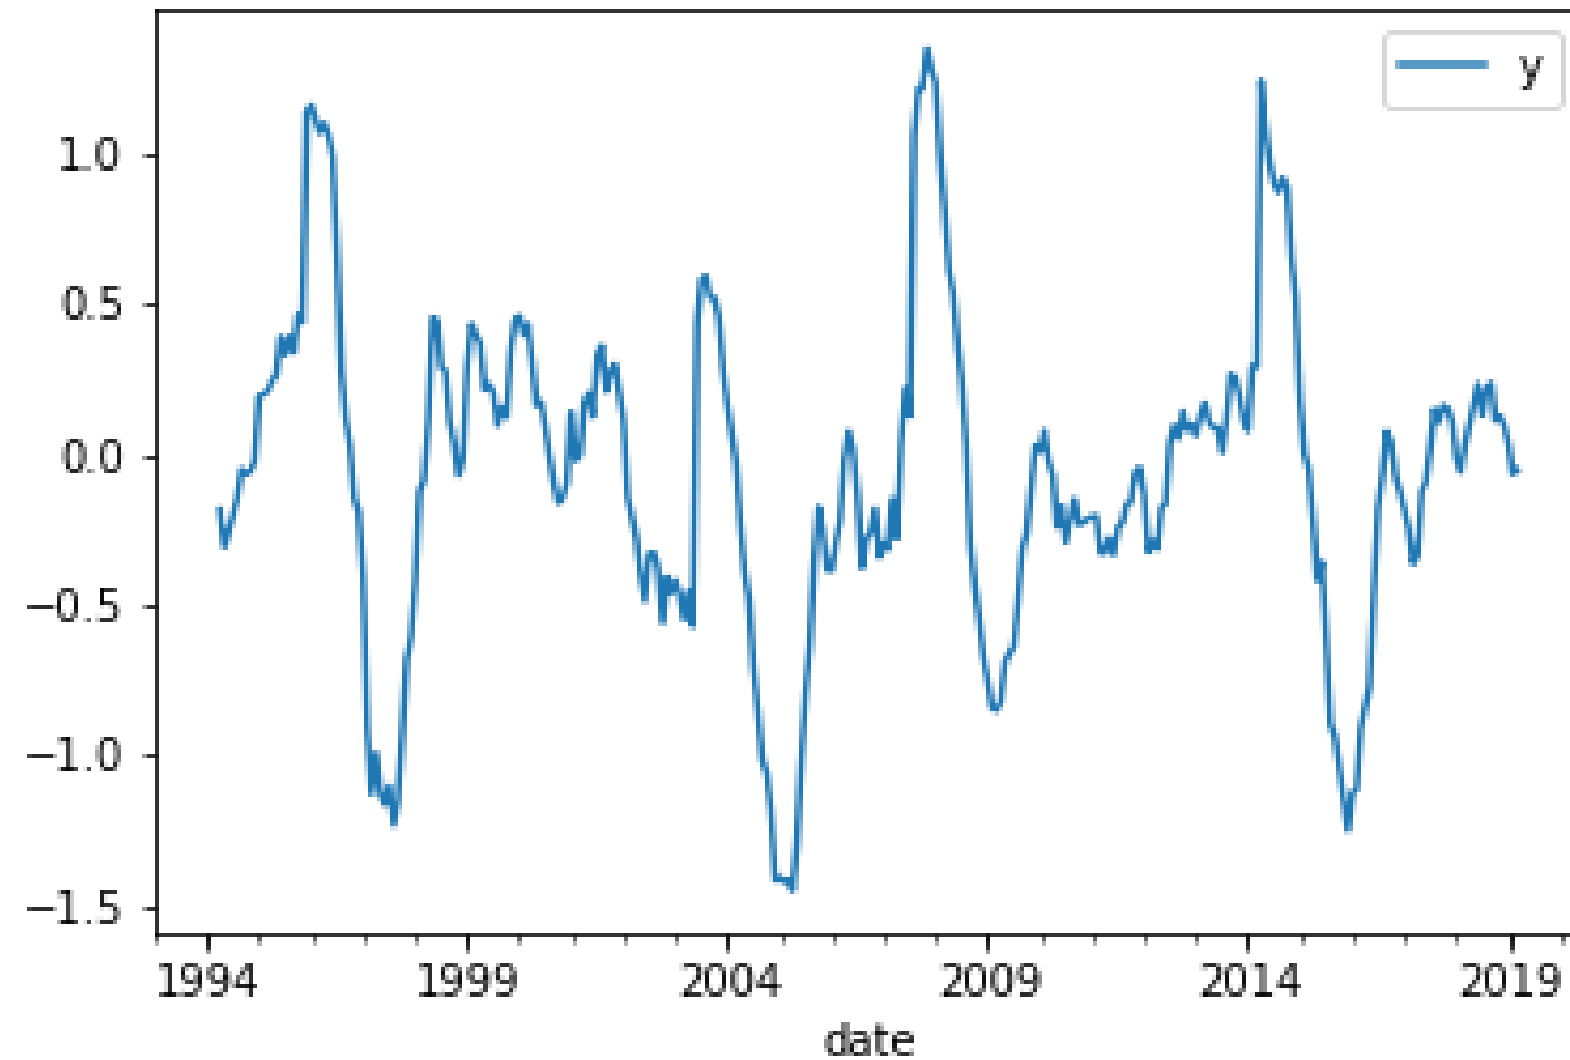
```
fig, ax = plt.subplots()
df.plot(ax=ax)
plt.show()
```



Seasonality



Cyclicality




White noise

White noise series has uncorrelated values


- Heads, heads, heads, tails, heads, tails, ...
- 0.1, -0.3, 0.8, 0.4, -0.5, 0.9, ...

Stationarity

Stationary




Not stationary



- Trend stationary: Trend is zero


Stationarity

Stationary




- Trend stationary: Trend is zero
- Variance is constant

Not stationary




Stationarity

Stationary



Not stationary



- Trend stationary: Trend is zero
- Variance is constant
- Autocorrelation is constant

Train-test split

```
# Train data - all data up to 30th March 2019
df_train = df.loc[:'2019-03-30']

# Test data - all data after 30th March 2019
df_test = df.loc['2019-03-30':]
```

Let's Practice!

FORECASTING USING ARIMA MODELS IN PYTHON

Making time series stationary

FORECASTING USING ARIMA MODELS IN PYTHON



James Fulton

Climate informatics researcher

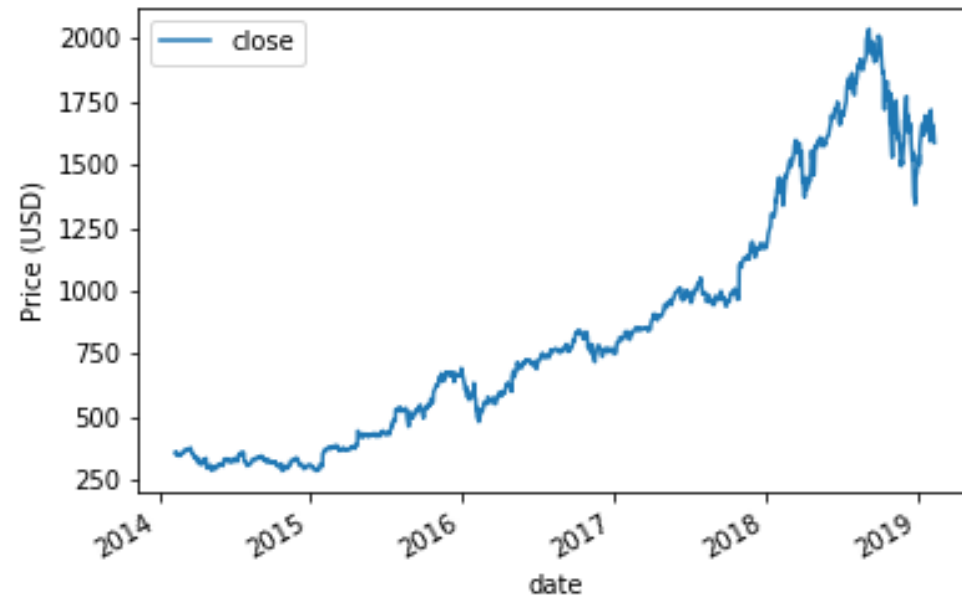
Overview

- Statistical tests for stationarity
- Making a dataset stationary

The augmented Dicky-Fuller test

- Tests for trend non-stationarity
- Null hypothesis is time series is non-stationary

Applying the adfuller test



```
from statsmodels.tsa.stattools import adfuller
```

```
results = adfuller(df['close'])
```

Interpreting the test result

```
print(results)
```

```
(-1.34, 0.60, 23, 1235, {'1%': -3.435, '5%': -2.913, '10%': -2.568}, 10782.87)
```

- 0th element is test statistic (-1.34)
 - More negative means more likely to be stationary
- 1st element is p-value: (0.60)
 - If p-value is small → reject null hypothesis. Reject non-stationary.
- 4th element is the critical test statistics

Interpreting the test result

```
print(results)
```

```
(-1.34, 0.60, 23, 1235, {'1%': -3.435, '5%': -2.863, '10%': -2.568}, 10782.87)
```


- 0th element is test statistic (-1.34)
 - More negative means more likely to be stationary
- 1st element is p-value: (0.60)
 - If p-value is small → reject null hypothesis. Reject non-stationary.
- 4th element is the critical test statistics

¹ <https://www.statsmodels.org/dev/generated/statsmodels.tsa.stattools.adfuller.html>

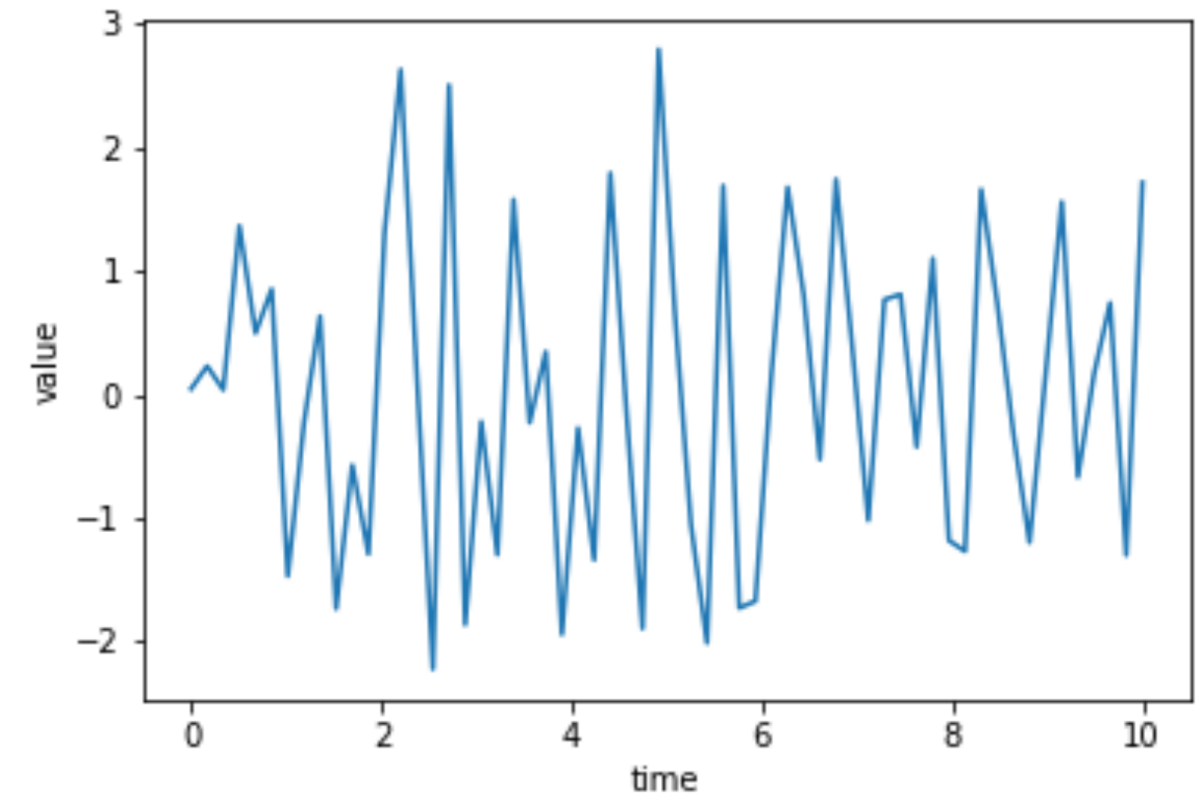

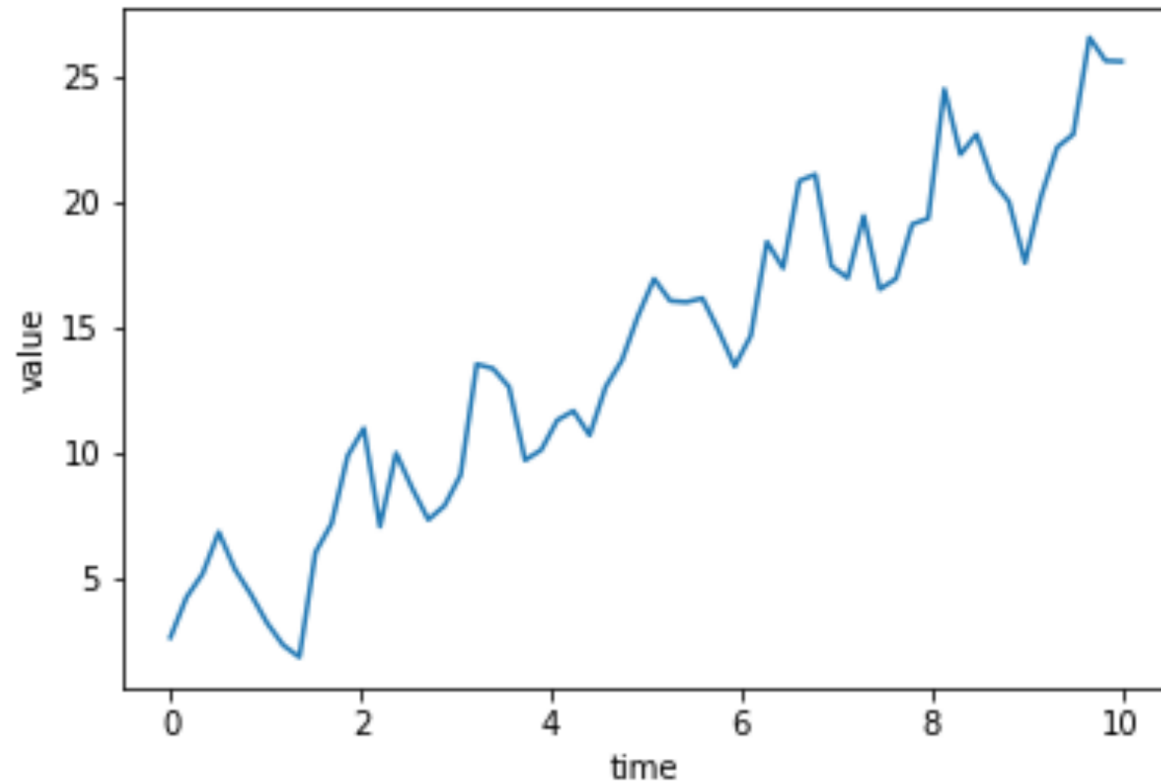
The value of plotting

- Plotting time series can stop you making wrong assumptions

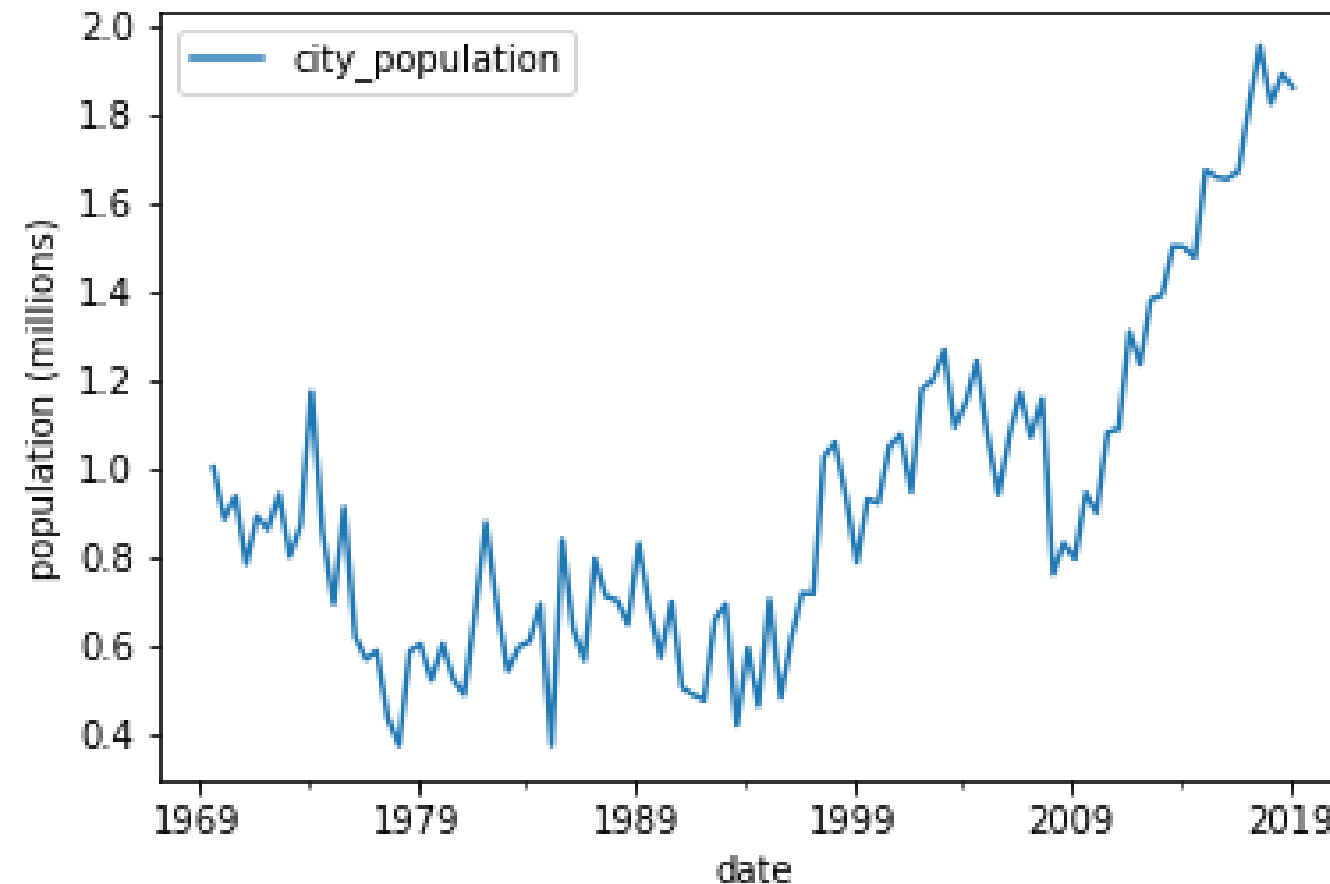
The value of plotting



Making a time series stationary



Taking the difference



Difference: $\Delta y_t = y_t - y_{t-1}$

Taking the difference

```
df_stationary = df.diff()
```

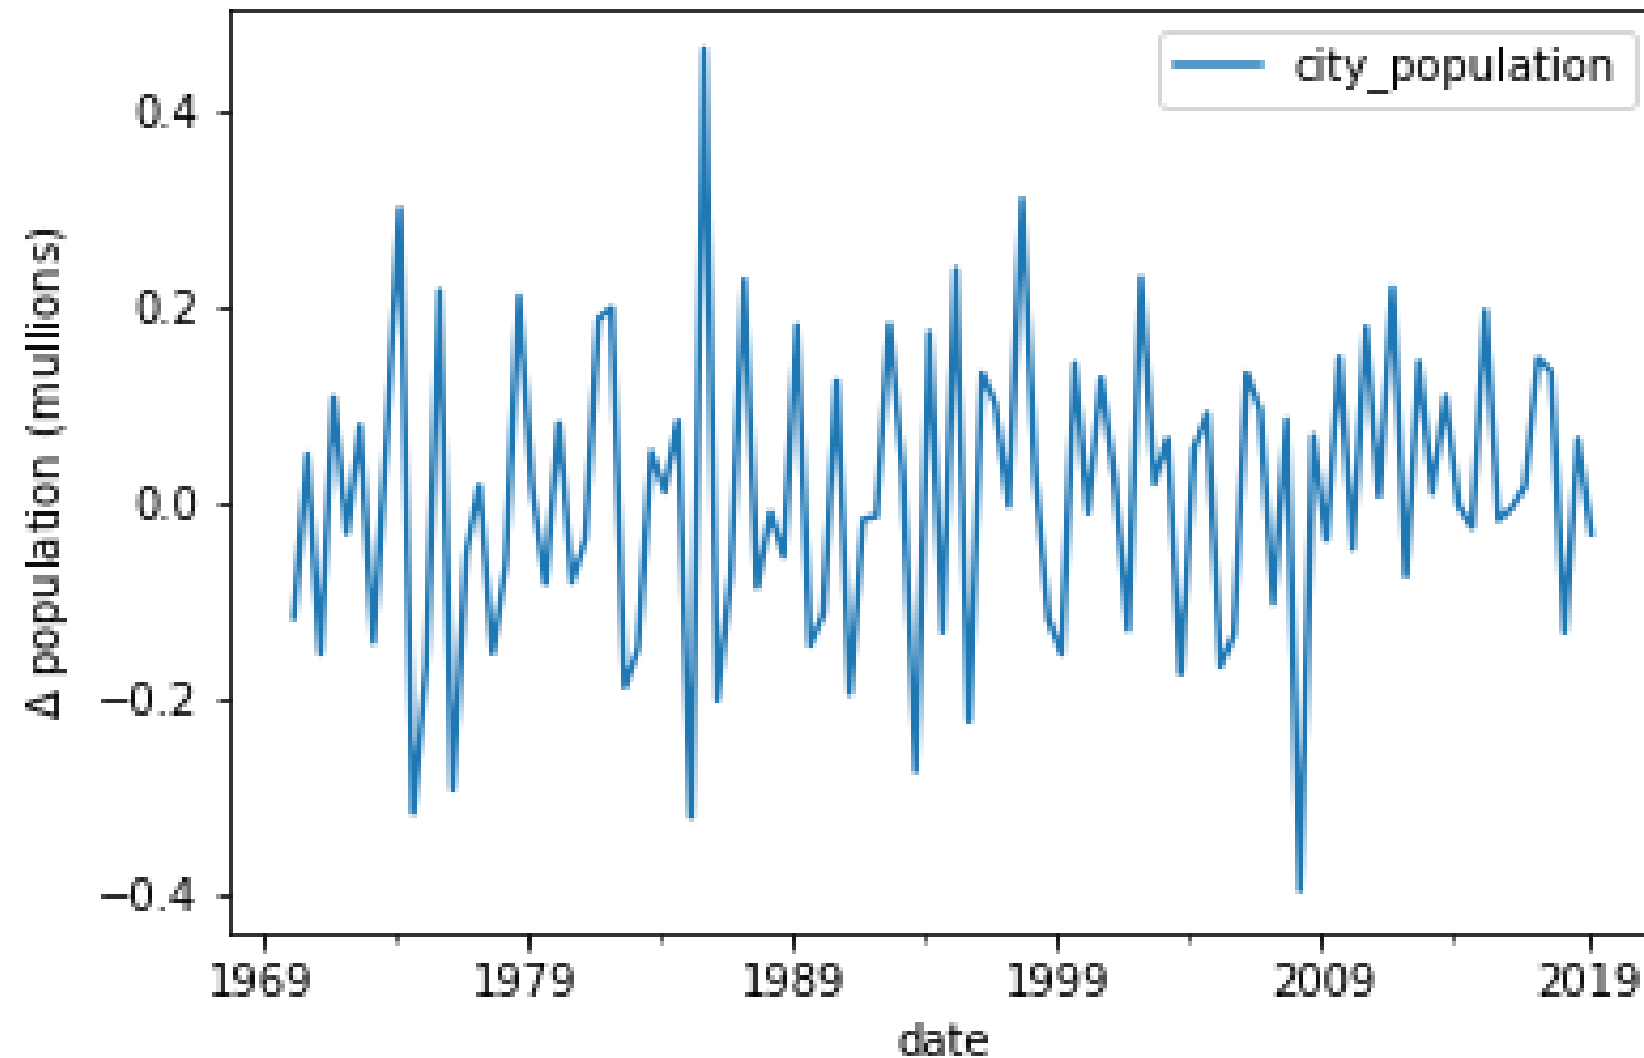
```
      city_population
date
1969-09-30         NaN
1970-03-31    -0.116156
1970-09-30     0.050850
1971-03-31    -0.153261
1971-09-30     0.108389
```


Taking the difference

```
df_stationary = df.diff().dropna()
```

	city_population
date	
1970-03-31	-0.116156
1970-09-30	0.050850
1971-03-31	-0.153261
1971-09-30	0.108389
1972-03-31	-0.029569

Taking the difference



Other transforms

Examples of other transforms

- Take the log
 - `np.log(df)`
- Take the square root
 - `np.sqrt(df)`
- Take the proportional change
 - `df.shift(1)/df`

Let's practice!

FORECASTING USING ARIMA MODELS IN PYTHON

Intro to AR, MA and ARMA models

FORECASTING USING ARIMA MODELS IN PYTHON



James Fulton


Climate informatics researcher

AR models

Autoregressive (AR) model

AR(1) model:

$$y_t = a_1 y_{t-1} + \epsilon_t$$



AR models

Autoregressive (AR) model

AR(1) model :

$$y_t = a_1 y_{t-1} + \epsilon_t$$

AR(2) model :

$$y_t = a_1 y_{t-1} + a_2 y_{t-2} + \epsilon_t$$

AR(p) model :

$$y_t = a_1 y_{t-1} + a_2 y_{t-2} + \dots + a_p y_{t-p} + \epsilon_t$$

MA models

Moving average (MA) model

MA(1) model :

$$y_t = m_1 \epsilon_{t-1} + \epsilon_t$$

MA(2) model :

$$y_t = m_1 \epsilon_{t-1} + m_2 \epsilon_{t-2} + \epsilon_t$$

MA(q) model :

$$y_t = m_1 \epsilon_{t-1} + m_2 \epsilon_{t-2} + \dots + m_q \epsilon_{t-q} + \epsilon_t$$

ARMA models

Autoregressive moving-average (ARMA) model

- ARMA = AR + MA

ARMA(1,1) model:

$$y_t = a_1 y_{t-1} + m_1 \epsilon_{t-1} + \epsilon_t$$

ARMA(p, q)

- p is order of AR part
- q is order of MA part

Creating ARMA data

$$y_t = a_1 y_{t-1} + m_1 \epsilon_{t-1} + \epsilon_t$$

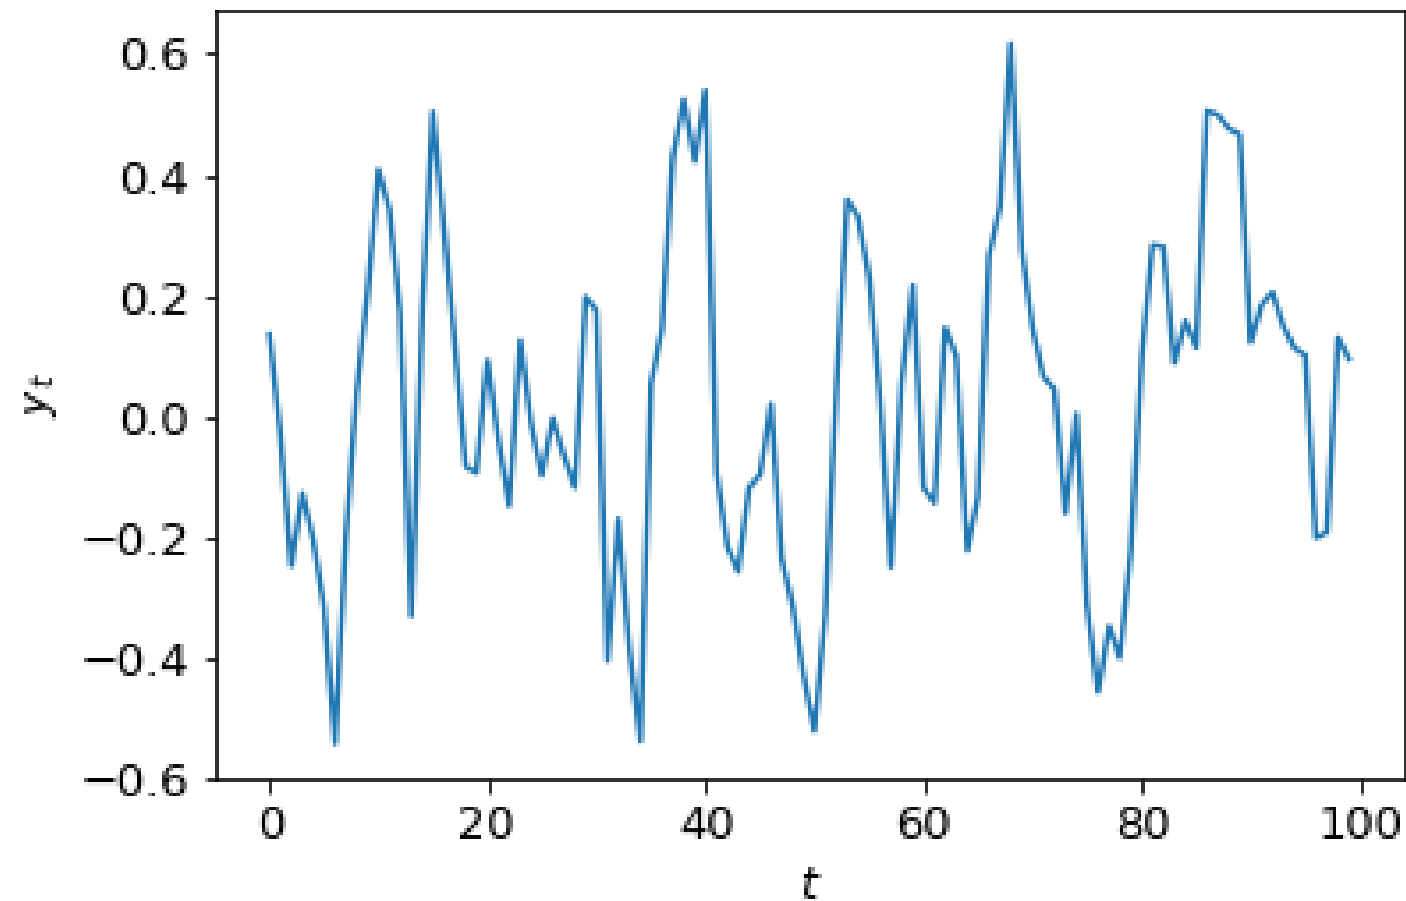
Creating ARMA data

$$y_t = 0.5y_{t-1} + 0.2\epsilon_{t-1} + \epsilon_t$$

```
from statsmodels.tsa.arima_process import arma_generate_sample  
  
ar_coefs = [1, -0.5]  
ma_coefs = [1, 0.2]  
  
y = arma_generate_sample(ar_coefs, ma_coefs, nsample=100, sigma=0.5)
```

Creating ARMA data

$$y_t = 0.5y_{t-1} + 0.2\epsilon_{t-1} + \epsilon_t$$



Fitting and ARMA model

```
from statsmodels.tsa.arima_model import ARMA

# Instantiate model object
model = ARMA(y, order=(1,1))

# Fit model
results = model.fit()
```

Let's practice!

FORECASTING USING ARIMA MODELS IN PYTHON