



MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

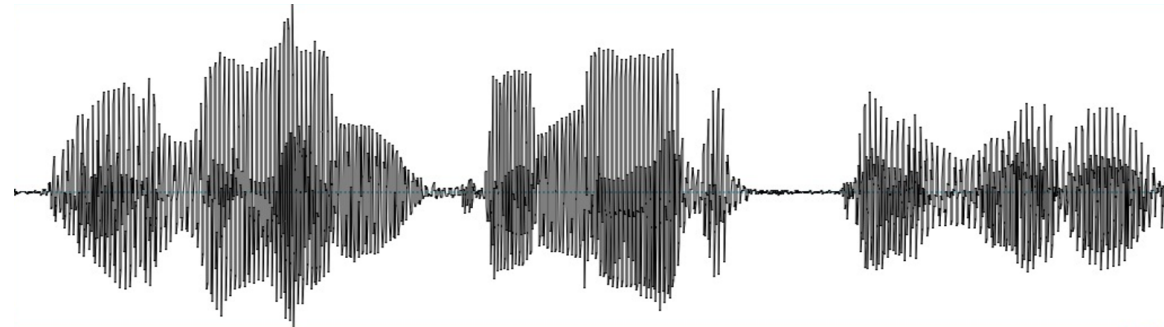
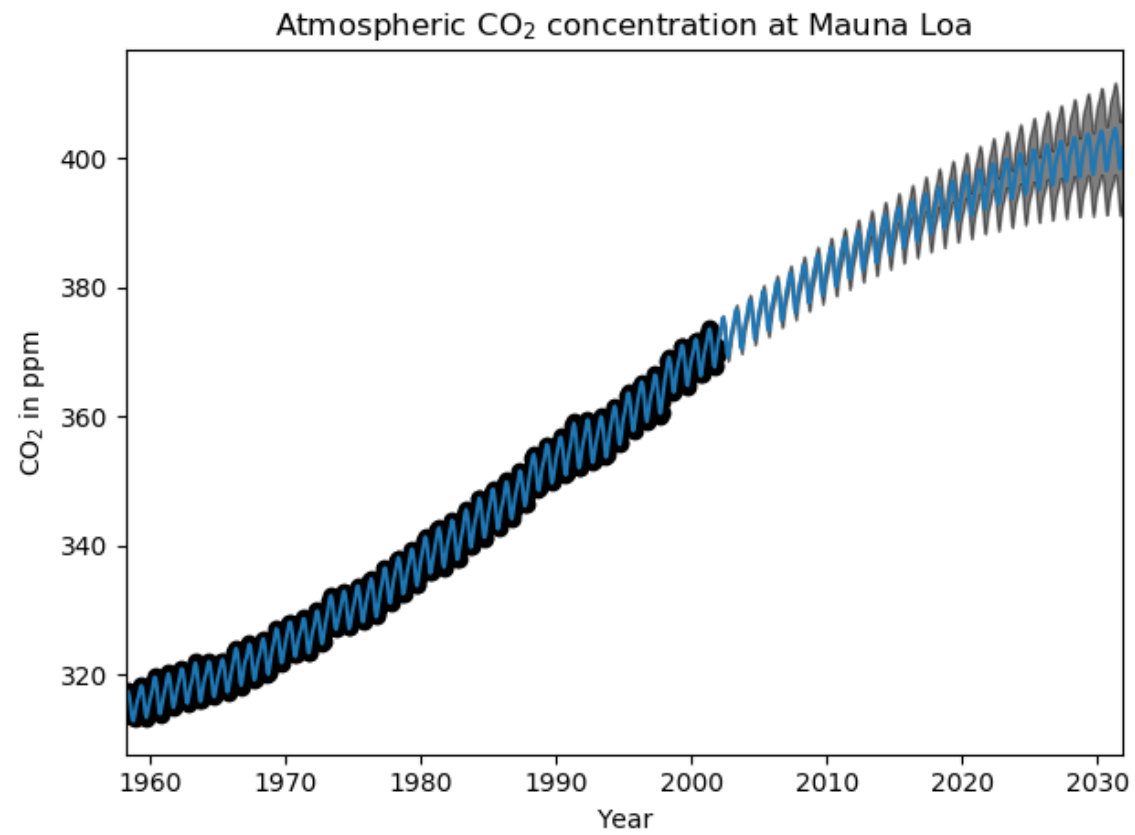
# **Timeseries kinds and applications**

**Chris Holdgraf**

Fellow, Berkeley Institute for Data Science

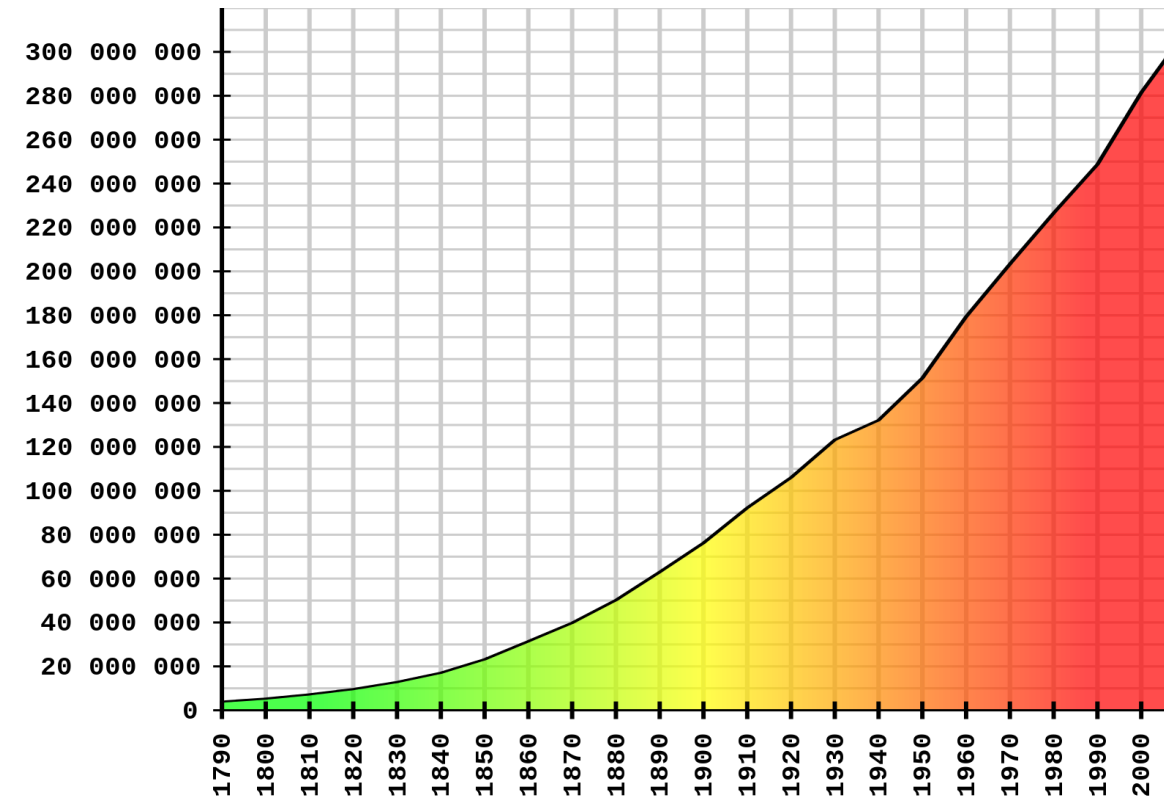
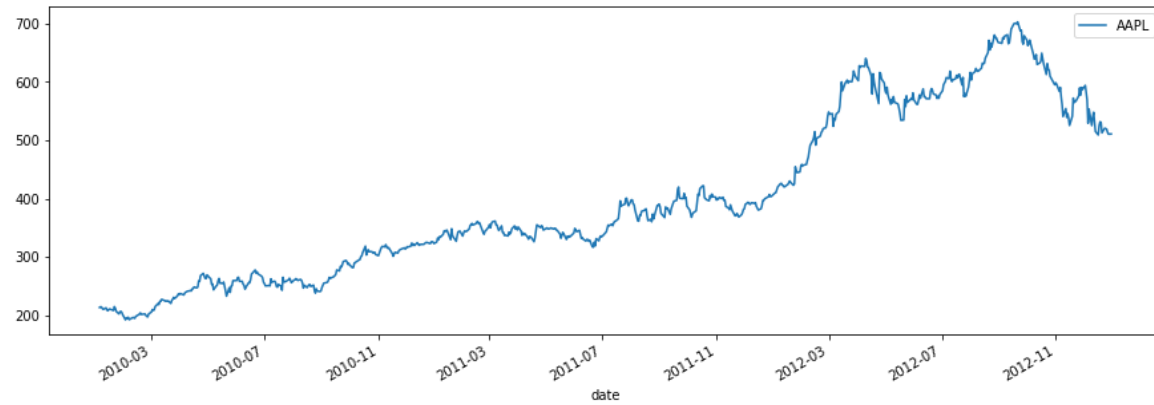


# Time Series





# Time Series





# What makes a time series?

Datapoint	Datapoint	Datapoint	Datapoint	Datapoint	Datapoint
1	34	12	54	76	40

Timepoint	Timepoint	Timepoint	Timepoint	Timepoint	Timepoint
2:00	2:01	2:02	2:03	2:04	2:05

Timepoint	Timepoint	Timepoint	Timepoint	Timepoint	Timepoint
Jan	Feb	March	April	May	Jun

Timepoint	Timepoint	Timepoint	Timepoint	Timepoint	Timepoint
1e-9	2e-9	3e-9	4e-9	5e-9	6e-9



# Reading in a time series with Pandas

```
import pandas as pd
import matplotlib.pyplot as plt
data = pd.read_csv('data.csv')
data.head()
```

	date	symbol	close	volume
0	2010-01-04	AAPL	214.009998	123432400.0
46	2010-01-05	AAPL	214.379993	150476200.0
92	2010-01-06	AAPL	210.969995	138040000.0
138	2010-01-07	AAPL	210.580000	119282800.0
184	2010-01-08	AAPL	211.980005	111902700.0

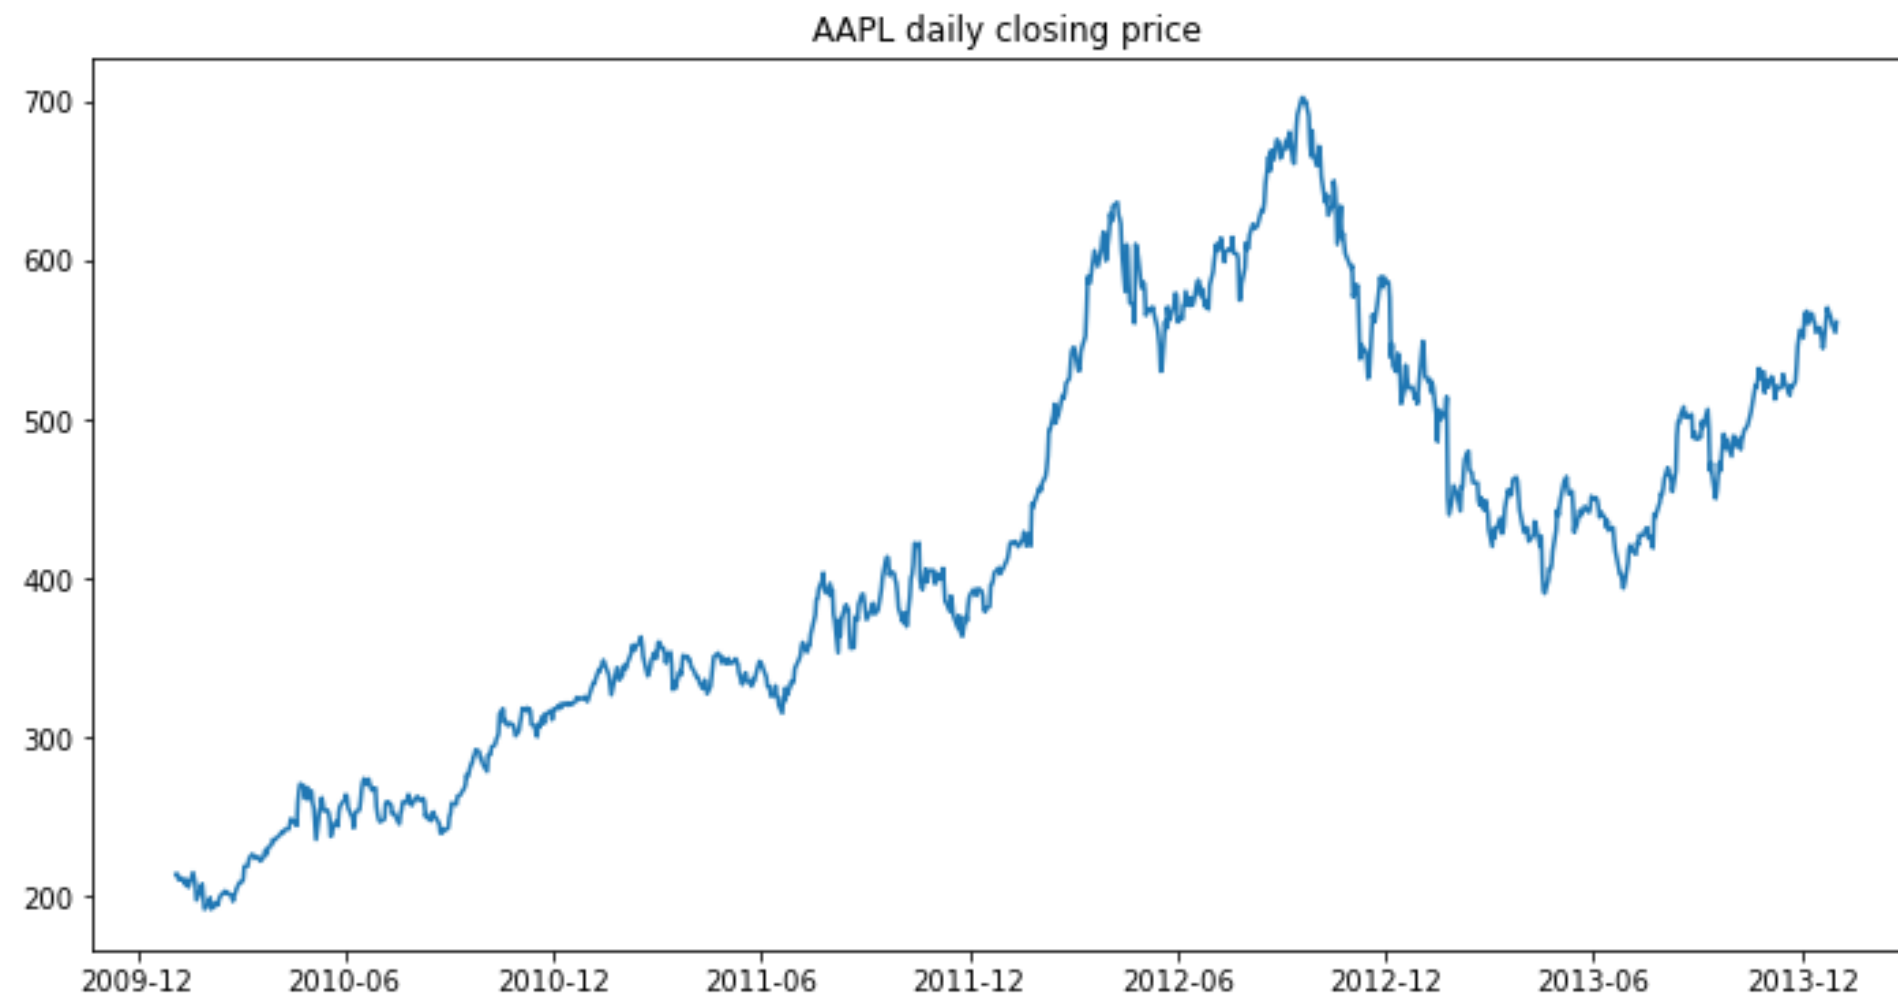


# Plotting a pandas timeseries

```
import matplotlib.pyplot as plt
fig, ax = plt.subplots(figsize=(12, 6))
data.plot('date', 'close', ax=ax)
ax.set(title="AAPL daily closing price")
```



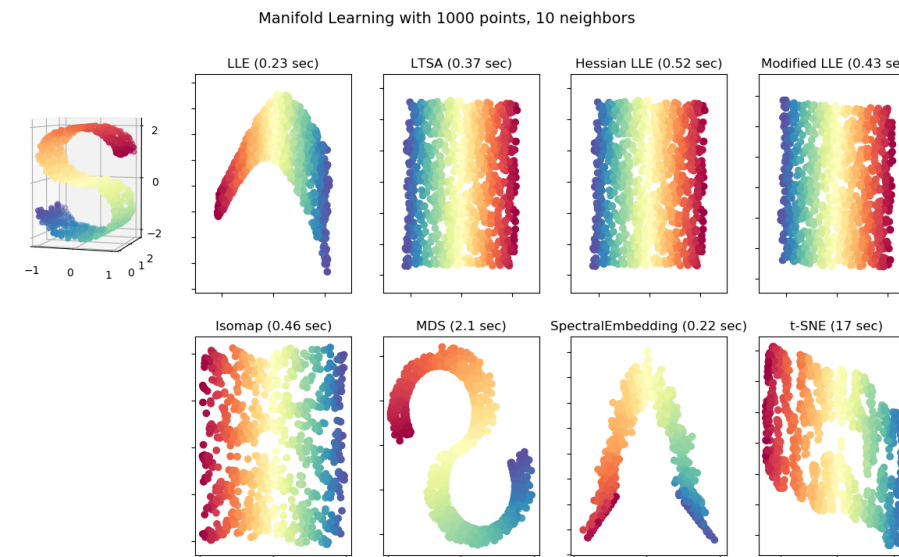
# A timeseries plot



# Why machine learning?

We can...

- Use really big data
- Use really complicated data






# Why machine learning?

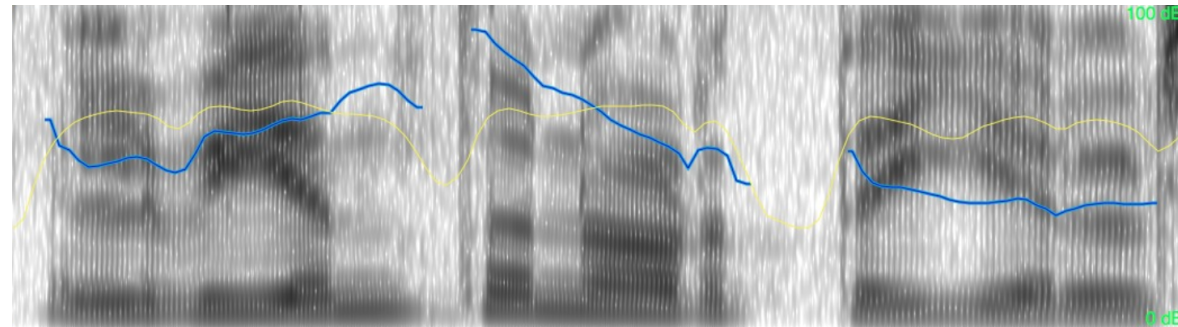
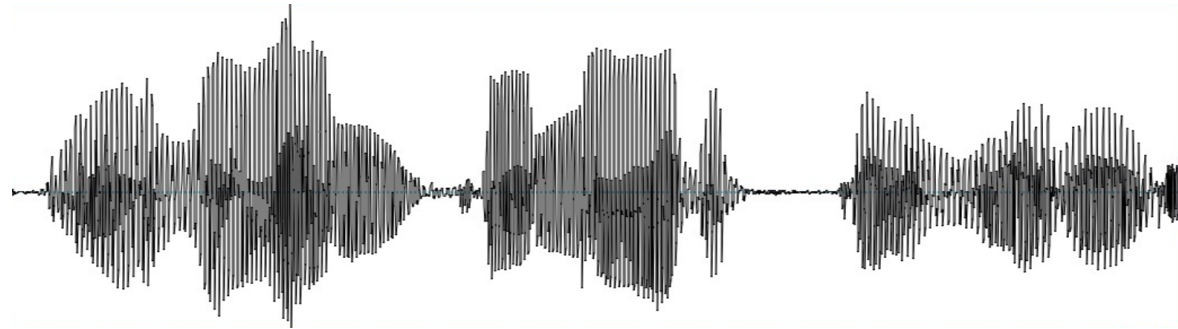
We can...

- Predict the future
- Automate this process





# Why combine these two?





# A machine learning pipeline

- Feature extraction
- Model fitting
- Prediction and validation



## MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

**Let's practice!**



MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

# Machine learning basics

Chris Holdgraf

Fellow, Berkeley Institute for Data Science



# Always begin by looking at your data

```
array.shape
(10, 5)

array[:3]
array([[ 0.735528,  1.00122818, -0.28315978],
       [-0.94478393,  0.18658748, -0.00241224],
       [-0.74822942, -1.46636618,  0.69835096]])
```



# Always begin by looking at your data

```
df.head()
```

	col1	col2	col3
0	0.735528	1.001228	-0.283160
1	-0.944784	0.186587	-0.002412
2	-0.748229	-1.466366	0.698351
3	1.038589	-0.171248	0.831457
4	-0.161904	0.003972	-0.321933



# Always visualize your data

**Make sure it looks the way you'd expect.**

```
# Using matplotlib
fig, ax = plt.subplots()
ax.plot(...)

# Using pandas
fig, ax = plt.subplots()
df.plot(..., ax=ax)
```





# Scikit-learn

**Scikit-learn** is the most popular machine learning library in Python

```
from sklearn.svm import LinearSVC
```



# Preparing data for scikit-learn

- `scikit-learn` expects a particular structure of data:

**(SAMPLES, FEATURES)**

- Make sure that your data is *at least two-dimensional*
- Make sure the first dimension is *samples*



# If your data is not shaped properly

- If the axes are swapped:

```
array.T.shape  
(10, 3)
```



# If your data is not shaped properly

- If we're missing an axis, use `.reshape()`:

```
array.shape  
(10,)  
  
array.reshape([-1, 1]).shape  
(10, 1)
```

- `-1` will automatically fill that axis with remaining values



# Fitting a model with scikit-learn

```
# Import a support vector classifier
from sklearn.svm import LinearSVC

# Instantiate this model
model = LinearSVC()

# Fit the model on some data
model.fit(X, y)
```

It is common for `y` to be of shape `(samples, 1)`



# Investigating the model

```
# There is one coefficient per input feature
model.coef_
array([[ 0.69417875, -0.5289162 ]])
```



# Predicting with a fit model

```
# Generate predictions  
predictions = model.predict(X_test)
```



## MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

**Let's practice**





MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

# Combining timeseries data with machine learning

Chris Holdgraf

Fellow, Berkeley Institute for Data Science



# Getting to know our data

- The datasets that we'll use in this course are all freely-available online
- There are many datasets available to download on the web, the ones we'll use come from Kaggle



# The Heartbeat Acoustic Data

- Many recordings of heart sounds from different patients
- Some had normally-functioning hearts, others had abnormalities
- Data comes in the form of audio files + labels for each file
- Can we find the "abnormal" heart beats?



# Loading auditory data

```
from glob import glob
files = glob('data/heartbeat-sounds/files/*.wav')

print(files)

['data/heartbeat-sounds/proc/files/murmur__201101051104.wav',
 ...
 'data/heartbeat-sounds/proc/files/murmur__201101051114.wav']
```



# Reading in auditory data

```
import librosa as lr
# `load` accepts a path to an audio file
audio, sfreq = lr.load('data/heartbeat-sounds/proc/files/murmur__201101051104.wav')

print(sfreq)
2205
```

In this case, the sampling frequency is 2205, meaning there are 2205 samples per second



# Inferring time from samples

- If we know the sampling rate of a timeseries, then we know the timestamp of each datapoint *relative to the first datapoint*
- Note: this assumes the sampling rate is fixed and no data points are lost



# Creating a time array (I)

- Create an array of indices, one for each sample, and divide by the sampling frequency

```
indices = np.arange(0, len(audio))  
time = indices / sfreq
```



# Creating a time array (II)

- Find the time stamp for the  $N-1$ th data point. Then use `linspace()` to interpolate from zero to that time

```
final_time = (len(audio) - 1) / sfreq  
time = np.linspace(0, final_time, sfreq)
```





# The New York Stock Exchange dataset

- This dataset consists of company stock values for 10 years
- Can we detect any patterns in historical records that allow us to predict the value of companies in the future?



# Looking at the data

```
data = pd.read_csv('path/to/data.csv')  
  
data.columns  
Index(['date', 'symbol', 'close', 'volume'], dtype='object')
```

```
data.head()
```

	date	symbol	close	volume
0	2010-01-04	AAPL	214.009998	123432400.0
1	2010-01-04	ABT	54.459951	10829000.0
2	2010-01-04	AIG	29.889999	7750900.0
3	2010-01-04	AMAT	14.300000	18615100.0
4	2010-01-04	ARNC	16.650013	11512100.0



# Timeseries with Pandas DataFrames

- We can investigate the object type of each column by accessing the `dtypes` attribute

```
df['date'].dtypes
```

```
0    object  
1    object  
2    object  
dtype: object
```



# Converting a column to a time series

- To ensure that a column within a DataFrame is treated as time series, use the `to_datetime()` function

```
df['date'] = pd.to_datetime(df['date'])

df['date']

0    2017-01-01
1    2017-01-02
2    2017-01-03
Name: date, dtype: datetime64[ns]
```



## MACHINE LEARNING FOR TIME SERIES DATA IN PYTHON

**Let's practice!**